

---

Research Interests    LLM and agent evaluation; AI auditing and accountability; abstention in autonomous agents; human–AI decision making; participatory evaluation; AI policy and governance.

## EDUCATION

2023–Now    **Brown University**, Providence, RI  
Ph.D. in Computer Science  
Adviser: Suresh Venkatasubramanian

2023–2025    **Brown University**, Providence, RI  
M.Sc. in Computer Science

2015–2020    **University of Ibadan**, Oyo, Nigeria  
B.Sc. in Computer Science

## RESEARCH EXPERIENCE

2023–Now    **Brown University**, Center for Technological Responsibility  
Graduate Fellow, advised by Suresh Venkatasubramanian  
Measurement frameworks and tooling for sociotechnical evaluation of LLMs and LLM-based agents: auditing workflows, abstention competence, and analysis of model behavior in deployment contexts. Evaluation protocols and datasets for robustness, multilingual performance, and reliability-oriented failure modes.

2022–2025    **Mozilla Foundation**, Open Source Audit Tooling (OAT) Project  
Research Assistant  
Research on AI auditing practice and audit tooling, resulting in two peer-reviewed publications and additional reports. Co-authored public comments to U.S. and E.U. agencies on AI accountability and standards.

2020–2023    **Brown University**, RISE Lab  
Research Assistant, working with C. Malik Boykin  
Research software and study design for work on user preferences over ML performance metrics and on uncertainty expression in decision-support contexts; participant-facing tutorials, data collection, and analysis.

2020–2023    **SeqHub Analytics LLC**  
Data Scientist  
Transformer-based document summarization; NLP features for a conversational assistant for educators.

## HONORS & AWARDS

|      |                                                                                     |
|------|-------------------------------------------------------------------------------------|
| 2026 | OpenMined Research Fellowship                                                       |
| 2026 | Best Paper Award, RLEval Workshop, ACM CAIS ★                                       |
| 2024 | Distinguished Paper Award, IEEE SaTML ★                                             |
| 2024 | National Artificial Intelligence Research Resource (NAIRR) Pilot                    |
| 2020 | Winner, Can You Hack the Future Ideathon                                            |
| 2020 | Winner, Best Problem Documentation, AI Commons Health Hackathon                     |
| 2019 | Nominated, Computer Science Department Programmer of the Year, University of Ibadan |

## SELECTED PAPERS

### *Preprints & Work in Progress*

|      |                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
|------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 2026 | <b>The AI Observatory: A Public Measure of Real-World AI Use</b><br>Shayne Longpre, Anka Reuel, Dayeon Ki, Cathy Mengying Fang, Megan Richards, Zhiping Zhang, Chuanyang Jin, Jenn Mickel, Cedric Deslandes Whitney, Ahmet Üstün, Niloofar Mireshghallah, <b>Victor Ojewale</b> , Ariel Lee, Alex (Sandy) Pentland, Tianshi Li, Yuntian Deng, Mykel Kochenderfer, Sara Hooker, Sanmi Koyejo<br>Under review (NeurIPS Datasets & Benchmarks 2026) |
|      | <b>Audit Trails for Accountability in Large Language Models</b><br><b>Victor Ojewale</b> , Harini Suresh, Suresh Venkatasubramanian<br>arXiv preprint                                                                                                                                                                                                                                                                                            |

### *Conference & Journal Publications (Peer-Reviewed)*

|      |                                                                                                                                                                                                                                        |
|------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 2026 | <b>Designing for Doubt: The Case for Informed Abstention in Autonomous Agents</b><br><b>Victor Ojewale</b> , Suresh Venkatasubramanian<br>AAAI/ACM Conference on AI, Ethics, and Society (AIES)                                        |
|      | <b>MonitrLLM: A Community-Centered Evaluation Infrastructure for Large Language Models</b><br><b>Victor Ojewale</b> , Ro Encarnación, Suresh Venkatasubramanian, Danaë Metaxa<br>AAAI/ACM Conference on AI, Ethics, and Society (AIES) |
|      | <b>Multilingual Functional Evaluation for Large Language Models</b><br><b>Victor Ojewale</b> , Inioluwa Deborah Raji, Suresh Venkatasubramanian<br>Findings of the Association for Computational Linguistics (ACL)                     |

**More Is Not Better: Visual Uncertainty Cues and the Fragility of Trust Calibration in LLM-Assisted Decision Making**

**Victor Ojewale**, Julia Ryan, Suresh Venkatasubramanian, C. Malik Boykin  
Computers in Human Behavior: Artificial Humans

2025 **Towards AI Accountability Infrastructure: Gaps and Opportunities in AI Audit Tooling**

**Victor Ojewale**, Ryan Steed, Briana Vecchione, Abeba Birhane, Inioluwa Deb Raji  
ACM Conference on Human Factors in Computing Systems (CHI)  
*Cited in the 2026 AI Action Summit International AI Safety Report*

2024 **AI Auditing: The Broken Bus on the Road to AI Accountability**  
Abeba Birhane, Ryan Steed, **Victor Ojewale**, Briana Vecchione, Inioluwa Deb Raji  
IEEE Conference on Secure and Trustworthy Machine Learning (SaTML)  
**Distinguished Paper Award** ★

*Workshop Papers & Extended Abstracts (Lightly Peer-Reviewed)*

2026 **What Benchmarks Don't Measure: The Case for Evaluating Abstention Competence in Autonomous Agents**

**Victor Ojewale**, Suresh Venkatasubramanian  
RLEval Workshop, ACM CAIS **Best Paper Award** ★

**Beyond Static Leaderboards: A Roadmap to Naturalistic, Functional Evaluation of LLMs**

**Victor Ojewale**, Suresh Venkatasubramanian  
Workshop on Language Models for Underserved Communities (LM4UC), AAAI

2025 **Testing LLMs in a Sandbox Isn't Responsible. Focusing on Community Uses and Needs Is.**

Michelle L. Ding, Jo Kavishe, **Victor Ojewale**, Suresh Venkatasubramanian  
Workshop on Socially Responsible Language Modelling Research (SoLaR), COLM  
*Opinion abstract*

*Public Comments & Regulatory Filings*

2024 **Comment on NIST-2023-0309**

Inioluwa Deb Raji, Abeba Birhane, Briana Vecchione, Ryan Steed, **Victor Ojewale**  
Request for Information (RFI) Related to NIST's Assignments Under Sections 4.1, 4.5 and 11 of the Executive Order Concerning Artificial Intelligence, National Institute of Standards and Technology, February 2

- 2023 **Comment on NTIA-2023-0005**  
 Inioluwa Deb Raji, Briana Vecchione, Abeba Birhane, Ryan Steed, **Victor Ojewale**  
 National Telecommunications and Information Administration, June 15  
*Cited in the NTIA AI Accountability Policy Report*
- Feedback from the Mozilla Open Source Audit Tooling Project**  
 Inioluwa Deb Raji, Briana Vecchione, Abeba Birhane, Ryan Steed, **Victor Ojewale**  
 Delegated Regulation on Data Access Provided for in the Digital Services Act  
 (F3423931), European Commission, May 31

## TALKS

- 2026 **Grounded LLM Evaluations**  
 Bryant University, MSDS 640 Data Science and AI Capstone
- 2025 **Technical AI Governance in Practice: What Tools Miss, and Where We Go Next**  
 Workshop on Technical AI Governance, ICML *Keynote*
- Towards AI Accountability Infrastructure: Gaps and Opportunities in AI Audit Tooling**  
 Ohio Data Ethics Working Group, Ohio State University
- 2024 **Grounded LLM Evaluations**  
 Brown University, CSCI 1951-Z Fairness in Automated Decision Making  
 NorthEast HCI Meeting
- 2023 **Navigating the Open-Source Algorithm Audit Tooling Landscape**  
 MozFest *Panel, co-facilitated with Abeba Birhane*

## SERVICE

### *Organizing Committee*

- 2025 Socials Co-Chair, ACM Conference on Fairness, Accountability, and Transparency

### *Program Committee & Reviewing*

- 2026 AAAI/ACM Conference on AI, Ethics, and Society (AIES)  
 NeurIPS
- 2025 AAAI/ACM Conference on AI, Ethics, and Society (AIES)  
 International Conference on Computational Linguistics (COLING)  
 Regulatable ML Workshop, NeurIPS
- 2024 Regulatable ML Workshop, NeurIPS  
 EvalEval Workshop, NeurIPS